University of Southampton

Faculty of Physical Sciences and Engineering
School of Electronics and Computer Science

Mercury Delay Line Emulation:
Reconstructing Memory for EDSAC

by

Joshua Tyler
August 29, 2017

A dissertation submitted in partial fulfilment of the degree of
MSc Embedded Systems

Abstract

Contents

List of Acronyms

AC alternating current.

ARM advanced RISC machine.

CPLD complex programmable logic device.

DC direct current.

EDSAC electronic delay storage automatic calculator.

ENIAC electronic numerical integrator and computer.

FIFO first in, first out.

FPGA field programmable gate array.

GBP gain-bandwidth product.

GPIO general purpose input/output.

HDL hardware description language.

IC integrated circuit.

LC logic cell.

LUT look-up table.

LVCMOS low voltage complementary metal oxide semiconductor.

PCB printed circuit board.

PLL phase locked loop.

RAM random access memory.

RC resistor-capacitor.

TQFP thin quad flat package.

VHDL VHSIC hardware definition language.

List of Figures

List of Tables

Chapter 1

INTRODUCTION

The National Museum of Computing is currently hosting a project to reconstruct a very early computer: electronic delay storage automatic calculator (EDSAC), the exact nature of the machine will be discussed in Chapter 2, however the recreation of its memory is non trivial and is the aim of this project.

[This segment is out of place] Creating a faithful reproduction of this system poses challenges, namely: the expense of mercury, the health and safety implications of using mercury in a museum environment, and the technical challenges of the precise machining necessary for the steel tubes. As a result of this the reconstruction project currently intends to use magnetorestrictive delay lines [4].

As discussed at length in [2], this method is non-ideal since it is: anachronistic for the time, dissimilar in appearance to the original delay lines, and dissimilar in terms of electrical interface to the original. For this reason it was decided to investigate the use of modern technology to emulate the original delay lines. This emulation is required to be indistinguishable from the original in terms of appearance and electrical interface. The design of this system is the goal of this project.

Chapter 2

TECHNOLOGY REVIEW

This Chapter presents an overview of EDSAC, and an overview of the relevant literature surrounding its the memory architecture.

The data derived will be used to derive a specification for the recreated memory solution, presented in Chapter 3. It should be ntod that much of the literature presented comes from original documentation produced by Maurice Wilkes, the man who led the original EDSAC project.

2.1 EDSAC overview

EDSAC, pictured with two of its creators in Figure 2.1 [5], was the first practical digital stored program computer. This means that it was the first practical computer able to accept a program from the user, store it in memory, and execute it on the fly. In contrast to this earlier computers, such as electronic numerical integrator and computer (ENIAC), hard-coded programs using switches, in the case of ENIAC using 3,600 ten 10-way switches [6]. The only digital stored program computer earlier than EDSAC wa the Manchester small-scale experimental machine. This machine was not intended for general purpose computation however, but rather for testing of a new type of memory [7].


Figure 2.1: W.Renwick and M.Wilkes standing with EDSAC [1]

2.2 EDSAC Memory Architecture

EDSAC ran its first program in May 1949. This posed a significant design challenge for the memory of the machine. Transistors were not commercially available at all until 1951 [8],

and valves, whilst available, were physically large, and were fairly expensive. This meant that creation of even a modest amount of storage would not have been feasible. ENIAC used valves, but it also had very little memory. This was not a problem for its intended application, but would have posed a problem for a general purpose computing platform, such as EDSAC [9, p.208].

The solution chosen for EDSAC's storage problem was delay line memory. This was common with other early computers, and works via a fairly simple mechanism. Given a medium able to delay a pulse train by a certain amount, memory can be created by feeding the output of that delay medium back into the input. If the delay time is tuned to be an integer multiple of the system clock frequency, the system is able to store a sequence of bits proportional in length to the delay time. This principle is illustrated in Figure 2.2.

Delay lines exist in various forms, from magnetorestrictive delay lines which function by twisting one end of a coil of wire, then waiting for the stress to propagate to the other end of the wire, to electric delay lines which provide much smaller delays by sending electrical impulses down a length of coaxial wire or a printed circuit board (PCB) microstrip trace.


Figure 2.2: Demonstration of the principle of delay line memory, adapted from [2]

EDSAC used acoustic delay lines in the form of steel tubes, approximately 1.5 m long, and full of mercury. Impulses were inserted into one end of the tube via a quartz transducer, and reach the other end of the tube after a delay of approximately 1 ms. Here a second quartz transducer converts the incoming acoustic impulse into an electronic impulse.

These delay lines are used in two ways:

1. Short tubes used for the results of calculations.
2. Batteries of longer tubes used as the main memory store. This is analogous to random access memory (RAM) in a modern computer.

EDSAC stores words in units of 36 pulses, which is referred to as a minor cycle. 34 pulses are used to store the magnitude of a number, one stores the sign, and one acts as a space between numbers. This system was chosen to allow storage of ten digit numbers.

The shorter tubes are sized to store only a single minor cycle, but the longer tubes are long enough to store 576 pulses (16 minor cycles). The batteries of each of these tubes each contained 16 tubes, and EDSAC originally had two batteries, allowing for a total storage capacity of $16 \times 16 \times 2 = 512$ numbers.

2.2.1 Timing

The memory in EDSAC uses a circulating bit rate of 500 kHz. This is made up of a 0.9 μs pulse, and a 1.1 μs gap for each bit (although some literature specifies a pulse of 0.9 μs and a

gap of 1.0 μs, implying a 526 kHz bit rate). The pulse is a burst of 13.5 MHz carrier frequency if the bit is a logical 1, or it is 0V if the bit is a logical 0.

It should be noted that the reconstruction will use a bit rate slightly faster or slower than 500 kHz, because 500 kHz is an international distress frequency, and given the large wiring looms in EDSAC, it would be quite easy for EDSAC to become an unintentional transmitter of this frequency. This would initially appear to oppose a faithful recreation of EDSAC, however one needs to consider the reaction of the mercury originally used to changes in temperature. Mercury was chosen because its acoustic delay does not vary much with temperature. The variation was still large enough however to cause cause the system to break over the normal temperature range of a laboratory environment. To combat this, the original clock frequency was adjusted with temperature such that the delay of each delay line was an integer multiple of the clock period. Later this system was deemed unsatisfactory and the coffins were enclosed in a temperature controlled environment.

To allow consistent operation of the machine, it is desirable that the recreated EDSAC does not emulate this variation of delay dependant upon temperature, but the delay of each line should be tunable around the nominal delay given by Equation 2.1.

$$D = \frac{1}{500\,\text{kHz}} \times 576 = 1.15\,\text{ms} \tag{2.1}$$

In the regeneration portion of the circuitry, the pulses are demodulated from the 13.5 MHz carrier and stretched to approximately 1.9 μs long, i.e. just long enough that each pulse fails to overlap it's neighbour [9, p.212].

This is critical because it means that the pulses that have passed through the delay line are effectively resynchronised to the system clock. If the system did not work in this way, then the delay lines in the the original EDSAC would never have worked. The propagation time of acoustic waves through mercury varies slightly with temperature, and the cumulative effect of this through all the delay lines of the store would mean that pulses would be unlikely to align with the system clock at the other end.

Despite this resynchronisation after each delay line, the original EDSAC suffered a great deal from the variation of delay with temperature. Originally the system clock was adjusted whenever the temperature of the delay lines drifted out of synchronization, but later on in development the delay lines were put inside a temperature controlled oven.

The regeneration system implies that the maximum acceptable skew of the delay line from its nominal delay is ±500 ns. This does not, however, take into account other factors such as the jitter and longer term drift of the system clock, as well the slew rates of the analogue circuitry. Whilst the demodulating pulse is lengthened to 1.9 μs, it is unlikely to be consistently at it's peak voltage for this time, and so the output pulse is likely to have a better shape if the delay line produces an output in the middle of this period. Therefore creating a system which is an order of magnitude better than the above calculation, i.e. a maximum deviation from the nominal value of 50 ns, seems reasonable.

2.2.2 Electrical

Electrically speaking, EDSAC originally drove the delay lines with a nominal voltage of 25 V peak to peak, through a 70 Ω terminated transmission line. The loss in the delay lines was 69 dB, leading to an output voltage of approximately 10 mV.

Despite this, the recreation project has discovered that the regeneration circuitry actually feed the delay lines with a decreasing voltage as the pulses propagate through the lines. The

Figure 2.3: EDSAC pulse timing

voltage starts off at approximately 35 V for a signal fed to a delay line at the start of the store, but is reduced to approximately 25 V peak for the tubes at the end of the store.

In addition to this, problems were experienced with amplifying the low signal level output by the delay lines. Because the current wire delay line solution has flexibility in its output voltage, currently the delayed signal is output at 100 mV peak.

The 75 Ω transmission lines are alternating current (AC) coupled to the main EDSAC chassis, and the shield of the coaxial cable is referenced directly to earth. This poses a small problem for the reconstruction, because modern health and safety regulations dictate that the chassis of EDSAC should be earthed, but the electronics themselves are powered from an isolated 0 V rail with earth leakage protection in place.

This wouldn't be a problem except for the fact that the coaxial shield is earth referenced, meaning the coaxial return current goes to earth, not the isolated 0 V rail. In order to combat this, several capacitors have been added between ground and the 0 V rail, in order to allow the AC current to pass, but still provide direct current (DC) isolation. This means that care will be necessary when implementing the circuitry to power the delay line, as it cannot draw any DC power from the line.

2.2.3 Mechanical

The store delay lines originally consisted of banks of approximately 165.5 cm long steel tubes. The tubes are then held in an array using machined end-plates. An illustration of this is shown in Figure 2.4.

The exact dimensions, aside from the length, of each store tube are not explicitly stated in the available literature. However the dimensions of the short tubes are stated, with each short tube having an outer diameter of 4.44 cm, and an inner diameter of 2.86 cm [9, p. 213]. For the purposes of this project, the long tubes are therefore assumed to have similar dimensions to this.

The main restriction these dimensions place on the project is that the electronics must fit inside a tube of the corrrect length and outer diameter, so that the reconstructed system can be indistinguishable from the original in terms of form.

Figure 2.4: A battery of mercury delay lines in EDSAC [3]

Chapter 3

SPECIFICATION

The research of Chapter 2 has led to the derivation of a specification for the delay line which will be produced. This specification is detailed in Table 3.1.

Table 3.1: Delay line specification

Item	Specification	Justification
1	**Must** be capable of producing a delayed copy of the EDSAC pulse train presented to it's input	This is the primary function of the device.
2	**Must** be powered from the input signal driven by EDSAC, with only minimal non-intrusive modifications made to EDSAC.	The goal of the device is to faithfully recreate the appearance and electrical interface of EDSAC, and thus large modifications such as power supply wires must be avoided.
3	**Must** be able to have a nominal delay of 1.15 ms, adjustable by at least $\pm 10\%$.	1.15 ms is the nominal delay of a long tube, as discussed in Section 2.2.1. An adjustable delay allows synchronisation with the system clock, with may vary.
4	**Must** have a maximum per burst deviation from the nominal delay of 50 ns.	This ensures that the delay line output will be able to synchronise with the clock of EDSAC. 50 ns is the maximum deviation derived in Section 2.2.1
5	**Must** be able to interface with input waveforms AC coupled bursts of 13.5 MHz carrier, with peak voltages in the range of 25 V to 35 V.	This is necessary to mimic the performance of the original delay line, 25 V to 35 V is the range derived in Section 2.2.2.
6	**Must** be able to have an adjustable nominal output voltage in the range of 10 mV to 100 mV (peak to peak), driving into 70 Ω.	An adjustable output voltage in this range allows compatibility with both the original electrical interface, and that used by the reconstruction effort, 10 mV to 100 mV is the range derived in Section 2.2.2.
7	**Must** be encapsulated in a metal tube of 4.44 cm outer diameter, and 165.5 cm length.	This diameter allows the design to have the same appearance of the main memory store tubes of the original EDSAC design, the width and diameter are discussed in 2.2.3
8	**Must** be accompanied by a testing device capable of emulating the signals produced by EDSAC.	This allows the delay line to be tested separately to the reconstruction project.

Chapter 4

DELAY LINE DESIGN

This Chapter describes the development of the delay line itself, covering the architectural choices made, as well as the detailed design, development and testing. The overall block diagram of the delay line is shown in Figure 4.1.

The delay itself is implemented in the digital domain, with the goal of accurately emulating the analogue domain of the original design. This development is detailed in Section 4.1. The surrounding analogue circuitry translates between the logic level signal signals of the digital circuit, and the input and output.

Figure placeholder

Figure 4.1: The overall architecture of the delay line

4.1   Digital Design

The goal of the digital design is to delay the signal on it's input by a certain amount. If the input signal were arbitrary, then the optimal solution would be to clock a 1 b wide first in, first out (FIFO) buffer with a sampling clock. The delay would then be given by Equation 4.1, where t_d represents the delay time, f_s sampling clock frequency, and N the depth of the buffer.

$$t_d = \frac{1}{f_s} \times N \qquad (4.1)$$

This solution can be simplified by consideration of the fact that the input signal is not arbitrary, the characteristics or the signal are known from the research detailed in Section 2.2.1. It is known that:

- The signal will consist of 0.9 μs pulses of 13.5 MHz tone.
- Each tone burst will be separated by 1.1 μs.
- Each delay line can store a maximum of 576 pulses.

Using these characteristics it can be seen that a digital delay line only needs to sample store the time at which the rising edge of an incoming pulse is received. Since the packet length and modulation frequency is fixed, this can be asserted on the line a fixed delay later.

4.1.1 Architecture selection

Various architectures could be used to implement the system described in the previous section. The three most obvious methods of implementation being

- A microcontroller design.
- A discrete logic design
- A field programmable gate array (FPGA) design

Microcontrollers have the advantage of being comparatively cheap and readily available, in addition they typically have more than enough RAM available to implement the memory to store the array of times at which pulses arrived.

The principle disadvantage is that microcontrollers inherently to their architecture process a single thread of data at once. This means that even a tight processing loop which samples the input would add a much larger amount of jitter to the input compared to a hardware solution with the same clock rate. Fortunately however modern microcontroller architectures, typically have a large number of peripherals embedded in the silicon, which can remove computation from the core. This combined with interrupts can provide an architecture with very predictable latency.

An example implementation is detailed in Figure 4.2. This details a system, based upon a typical microcontroller in the advanced RISC machine (ARM) Cortex-M0 family. A pin change interrupt interrupts the main execution thread, and branches to an interrupt handler. This interrupt handler reads the value of a free-running timer peripheral that counts up using the microcontroller master clock, adds a fixed delay value to it, and appends it to a queue of timer counts held in RAM. This queue contains the values of the counter which the output should be asserted at.

The main execution thread of the microcontroller configures the timer peripheral to trigger a second timer peripheral once its count equals the value on the top of the queue. This second counter is connected directly to an output general purpose input/output (GPIO) pin, and is clocked such that the output toggles at 13.5 MHz.

This system can therefore meet the requirements of the delay line, with a worst case jitter of one system clock period (since the Cortex-M0 family allows for deterministic latency interrupts **[cite]**). Microcontrollers with timers capable of being configured as described above are readily available also **[cite STM32 reference manual here]** .


Figure 4.2: Proposed architecture for a microcontroller based system

The other two alternatives, a discrete logic system, and a FPGA based system would be similar in architecture, but differ in implementation, with the FPGA based system implementing the function using the programmable fabric of a FPGA, whereas the discrete implementation would use individual integrated circuits (ICs). A block diagram of the proposed architecture is shown in Figure 4.3.

The concept of this architecture is similar to that of the microcontroller based system. There is a free running counter inside the FPGA, when a pulse is detected on the input, the value of the counter, added to a fixed delay is saved into a FIFO. A second hardware module outputs a pulse train whenever the value of the counter matches the value on the top of the FIFO.

The difference between the hardware implementation and the microcontroller implementation is that there isn't a processor core to set up the hardware blocks, instead each hardware block is designed to perform the correct function, and interacts with the other modules using logic signals. In addition there is no need for an external modulator, as a harware block can be created to output the modulated 13.5 MHz signal.

The FIFO is the centre of this system. It is set to be the width of the free-running counter. At its input is the current value of the counter, added to a fixed constant that represents the required delay. This value is saved into the FIFO when the rising edge detector produces a pulse on its output.

The rising edge detector is a fairly simple hardware block that outputs a pulse for a single clock cycle when it detects a rising edge on its input. It then times out for a fixed interval, to avoid triggering on the remaining pulses in the same packet, and resets.

At the output of the FIFO feeds into a comparator. This comparator compares the value on the top of the FIFO with the current value of the counter, and triggers the pulse generator when the two values are equal.

The pulse generator outputs a fixed length burst of 13.5 MHz tone whenever it is triggered. This could be implemented simply by choosing a clock frequency that is 2^n times greater than 13.5 MHz. Thus the carrier frequency can be generated by a counter clocked from the input clock.
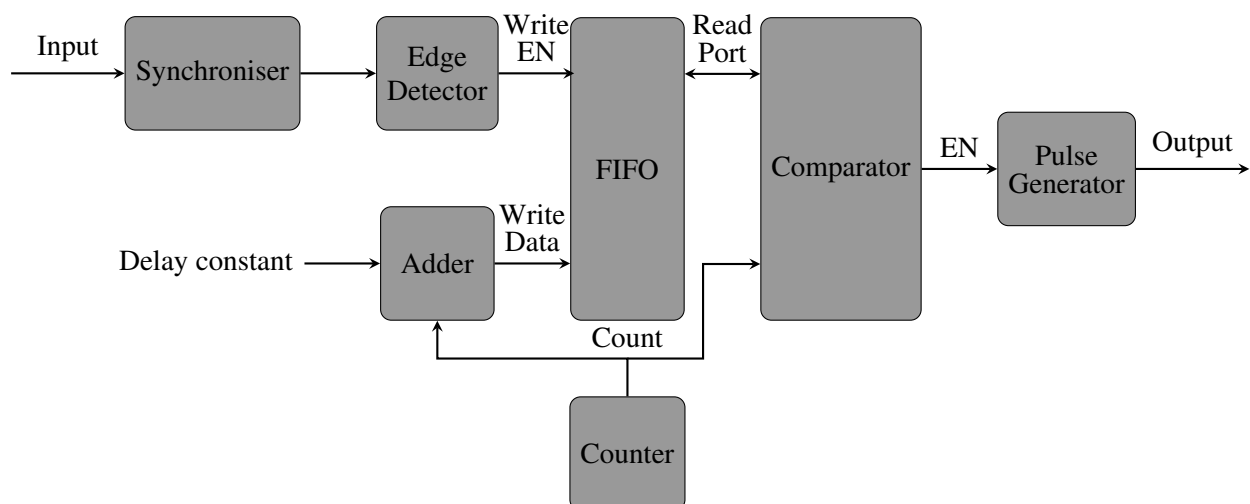

Figure 4.3: Proposed architecture for a FPGA or discrete system

[Could add power in here]

The advantage of the microcontroller system over a FPGA or discrete solution is twofold. Firstly the simplicity of the physical circuit necessary. A microcontroller based design would

require a IC for the microcontroller itself, and an external crystal oscillator (although some microcontrollers do have a reduced precision internal resistor-capacitor (RC) oscillator, removing the need even for this), and are generally powered from a single 3.3 V supply. This is in contrast to FPGAs which typically require more than one supply rail, with extensive decoupling, and a discrete solution which would require many ICs.

Secondly a microcontroller design would likely be slightly cheaper than a discrete logic design, as low end microcontrollers are typically only slightly more expensive than individual logic ICs, and are much less costly than typical FPGAs.

Despite these advantages, a hardware solution does have the advantage of elegance. While it has been demonstrated above that a microcontroller solution could achieve single cycle jitter – the same as a hardware solution. It is a lot more difficult to implement and verify, given the fact that configuration of many hardware peripherals are required.

In addition, the cost advantage of the microcontroller solution may be smaller than anticipated. This is because the simplicity of the hardware solution would only require a very small FPGA, which may be competitive in price to a microcontroller. Alternatively a complex programmable logic device (CPLD) with an external RAM IC could be used. A FPGA solution would be preferred to a discrete implementation, due to the ease of testing, and reconfiguring the hardware if the requirements change.

Therefore on balance it has been decided to proceed with a FPGA based solution.

FPGA Selection

The primary requirements for the FPGA are as described in Table 4.1.

Table 4.1: FPGA Requirements

Number	Requirement	Justification
1	**Must** have a moderately low cost	Using the IC for every delay line in the store must not be cost prohibitive.
2	**Must** have enough logic elements and block RAM to implement a single long delay line	This is the proposed function of the delay line
3	**Must** have fabric capable of being clocked fast enough to implement the design with a clock rate of [xx]	It was decided at [xx] that this is the required sampling speed.
4	**Should** have a development board available with a width less than [xx]	As described in Section [xx], [xx] is the internal diameter of the delay line tube, it would be ideal if the development board could fit inside the tube, to save a custom PCB being designed.
5	**Should** have a phase locked loop (PLL) to enable to fabric clock to be generated from a crystal oscillator	An internal PLL is ideal, but an external PLL could also be used.

At the time of writing the least expensive FPGA available from component suppliers Farnell, is the a 1280 logic cell variant of the iCE40 FPGA family, produced by Lattice [10]. This

is £5.10 in single quantity, which is low cost enough that it could reasonably be used to replace all of the delay lines, thus meeting specification point 1.

This FPGA has 1280 logic cells (LCs), which should be plenty to implement the simple design, given that each LC in this architecture consists of a four input look-up table (LUT), and a flip-flop [11, p.2-2]. The FPGA also has 64 kb of block RAM. This meets specification point 2, when one considers that is enough to store a clock sample for each of the 576 possible pulses, even if a 64 b clock width was used, as demonstrated by Equation 4.2, where S represents the size of memory required. In addition to this, there is one PLL available on the FPGA die, meeting specification point 5.

$$S = 64\,\text{b} \times 576 = 36\,\text{kb} \tag{4.2}$$

It is hard to estimate how fast a design will be able to operate in a FPGA without synthesising it, and running timing analysis. However, despite this we can estimate that the FPGA will easily meet timing at **[xx]** , thus meeting specification point 3, given that the register-to-register performance of the fabric is as good as 403 MHz for a dual-port RAM, 305 MHz for a 16:1 multiplexer, and 105 MHz for a 64 b counter.

In addition to this, a development board is available which is very narrow [12]. The exact dimensions are not provided, but it is barely wider than the 22 mm thin quad flat package (TQFP) of the FPGA itself [12, p.2], meaning it is highly likely to fit in the **[xx]** diameter tube, thus meeting specification point 4.

Based upon the fact that it meets all of the specification points, it has been decided to implement the design using the Lattice iCEstick evaluation board, with the possibility of moving to a custom PCB if many instances of the design were required.

4.1.2 HDL Design

Uniquely, the Lattice iCE40 family of FPGAs has an open-source toolchain available, named Project IceStorm which can be used as an alternative to the toolchain provided by Lattice [13]. Project IceStorm synthesises Verilog natively, and Lattice's iCEcube2 toolchain synthesises both Verilog and VHSIC hardware definition language (VHDL) [14, p.10]. Therefore in order to maintain compatibility with both toolchains, the design will be written using Verilog. Both toolchains were trialled, and the codebase is compatible with both, however Project IceStorm was used in the end as it is the more user-friendly toolchain.

The design was implemented using the structure of Figure 4.3, with the rising edge detector, FIFO, comparator, and pulse generator implemented as individual Verilog modules.
[Could talk about how the parameters are determined here]

Rising Edge Detector

The rising edge detector is implemented as the state machine of Figure 4.4. In the wait state, the input is sampled on every clock edge, when it is true, the state machine transitions to the assert state, where the output is asserted, until the state machine unconditionally branches to the timeout state. In the timeout state a counter is incremented until it reaches the required limit. At this point the state machine branches back to the wait state.
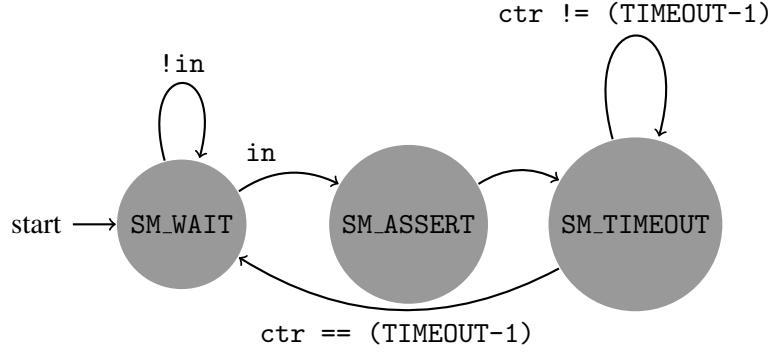
Figure 4.4: Rising Edge Detector State Machine

FIFO

The FIFO is implemented by building logic constructs around a dual port RAM to handle addressing of the read and write ports.

The main part of the FIFO is the read and write addresses. These addresses act as pointers to the current location of the read/write word in RAM. Each address is incremented on when the FIFO it read from/written to, and the address loops around to the start when it reaches the end of the RAM buffer. This 'round robin' approach means that the FIFO can be read from and written to an unlimited amount of times, so long as the total number of words stored is not greater than the depth of the buffer.

In addition to the read/write address pointers, there is a counter which keeps track of the total number of words stored in the buffer. This counter is decremented if a read is requested, and incremented if a write is requested (its value does not change if both a read and a write is requested at the same time, or neither a read or write is requested). This counter is used to generate empty and full signals so that the surrounding logic knows the state of the counter.

Comparator

The comparator is implemented as a state machine that requests data from the FIFO and asserts the output when the counter matches the data word, as illustrated by the state transition diagram of Figure 4.5.

`empty` is the `empty` signal from the FIFO, and the FIFO read request signal is true when the state machine is in the request state.

`count` is the value of the system counter, and `data_in` is the data read from the FIFO. The output trigger is true when the state machine is in the `SM_ASSERT` state.

4.1.3 HDL Verification

In order to verify the hardware description language (HDL) design, a SystemVerilog testbench was written for each module, and these modules were then simulated using Mentor Graphics ModelSim. These testbenches provide stimulus to verify that each block works as anticipated before integration with the system as a whole.

The focus of the verification is, however, on the system level testing, which tests the delay line as a whole. This testbench generates pulse trains representing random numbers in the same format as EDSAC. The number is then transmitted as a series of modulated pulses in the same way as EDSAC, and each pulse burst is added to a queue with the simulation time when the
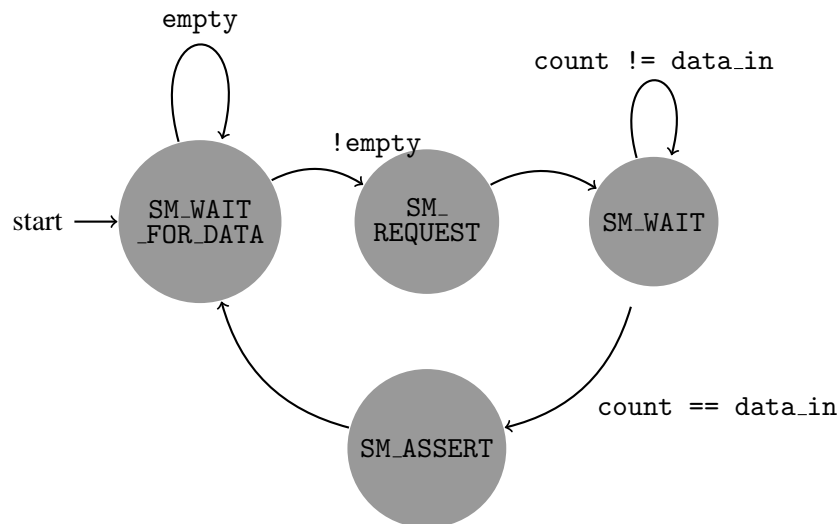
Figure 4.5: Comparator State Machine

pulse is expected to arrive. The testbench then checks the queue on each time step to verify that the system produces output bursts of the correct frequency at the correct time. Figure 4.6 shows the output of this testbench.


Figure placeholder

Figure 4.6: Delay line testbench output

4.2   Analogue Design

The FPGA communicates using 3.3 V low voltage complementary metal oxide semiconductor (LVCMOS) inputs and outputs. This is a great difference from the 25 V to 35 V input, and 10 mV to 100 mV output range. In addition to this, the input and outputs are AC coupled and so require a voltage symmetrical about 0 V.

4.2.1   Input

The input voltage needs to be terminated with 68 Ω, the closest E12 preferred resistor value to 70 Ω and the value which is used by the remainder of the reconstruction project to terminate

the transmission lines, in addition to this the power system imposes a requirement to remove the 50 Hz power signal discussed in Chapter 5.
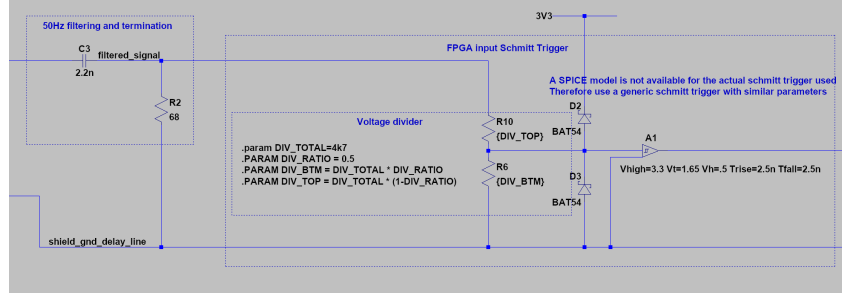

Figure 4.7: Delay line input schematic **[replace with better schematic]**

The schematic designed to interface the input signal to the FPGA is shown in Figure 4.7.

The 2.2 nF capacitor and 68 Ω resistor act as a low pass filter with a break frequency of 1.1 MHz, as derived in Equation 4.3. This correctly terminate the incoming 13.5 MHz bursts, as well as presenting a high impedance to the 50 Hz power supply signal.

$$f_0 = \frac{1}{2\pi RC} = \frac{1}{2\pi \times 68\,\Omega \times 2.2\,\text{nF}} = 1.1\,\text{MHz} \tag{4.3}$$

The next stage of the input circuitry is a 4.7 kΩ potentiometer. This attenuates the input signal by an adjustable amount to provide a logic level input signal for the Schmitt trigger buffer.

The value of the potentiometer is designed such that it provides negligible loading on the input signal compared to the 68 Ω termination resistor. Despite this the current through it for a 10 V input signal is $\frac{10\,\text{V}}{4.7\,\text{k}\Omega} = 2.1\,\text{mA}$, an order of magnitude greater than both the BAT54 diode leakage current (2 μA max [15, p.2]), and the Schmitt buffer input leakage current (5 μA max [16, p.4]).

The Schmitt trigger removes small glitches which may occur on the input so that they do not reach the FPGA input. it also helps to provide a defined LVCMOS output signal from the unclean input signal.

The BAT54 protection diodes clamp the input signal to the Schmitt buffer close to its supply rails. The are suitable for the task due to their fast recovery time, 5 ns and low forward voltage, 0.4 V at 10 mA [15, p.2]. In addition their continuous forward current rating, 200 mA max is far greater than what can be sourced through the 4.7 kΩ potentiometer.

4.2.2 Output

The output section of the delay line is required to drive the line with 0 V when no pulse is present, and a 10 mV to 100 mV peak-to-peak signal, centered about 0 V when no signal is present.

The biasing about 0 V is achieved by AC coupling the output of the amplification circuit. One problem which can occur from this is that the output of the FPGA will be 0 V when no signal is being output, but will oscillate between 0 V and 3.3 V when a pulse is being transmitted. This is a problem for the circuit because to correctly AC couple the signal, the voltage at the capacitor should be half way between the two extreme values, i.e. 1.65 V. This is illustrated in Figure 4.9.

In order to achieve this effect, the internal tri-state buffer capability of the FPGA is used. This is the purpose of the resistor network shown on the output of the FPGA. When the FPGA output is disabled, the two resistors between power rails pull the op-amp non-inverting input to $3.3\,\text{V} \times \frac{1\,\text{k}\Omega}{2\,\text{k}\Omega} = 1.65\,\text{V}$. When the FPGA drives its output to $3.3\,\text{V}$, the voltage at the op-amp is equal to $3.3\,\text{V} \times \frac{1\,\text{k}\Omega}{1.5\,\text{k}\Omega} = 2.2\,\text{V}$, and when the FPGA drives $0\,\text{V}$, the voltage at the op-amp will be $3.3\,\text{V} \times \frac{0.5\,\text{k}\Omega}{1.5\,\text{k}\Omega} = 1.1\,\text{V}$. This arrangement is used, rather than having the FPGA directly drive the non-inverting input of the op-amp in order to avoid driving the op-amp into saturation, as it is powered from the same $0\,\text{V}$ to $3.3\,\text{V}$ rails as the FPGA GPIO pins.

The first stage op amp is a unity gain buffer amplifier used to drive the output into a $220\,\Omega$ potentiometer. This potentiometer is used to attenuate the signal to the required level from the $1\,\text{V}$ peak to peak amplitude driven by the op-amp to the required level. This also attenuates the bias point away from half way between the voltage rails. For this reason, the AC component of the signal is re-biased onto $1.65\,\text{V}$ before buffering by the second amplifier.

The second amplifier is a second unity gain buffer, that critically has an output impedance of $68\,\Omega$ and AC couples the output signal in order to remove the DC bias.


Figure 4.8: Delay line output schematic **[replace with better schematic]**

For both amplifiers in the circuit LMH6639 amplifiers are used [17]. This amplifier is chosen because of several desirable properties. Firstly it has a gain-bandwidth product (GBP) of $190\,\text{MHz}$, more than what is necessary to buffer the $13.5\,\text{MHz}$ signal with unity gain. In addition to this, unlike many high GBP op-amps, the LMH6639 is a voltage feedback architecture simplifying design. Crucially for this application the amplifier is stable at a gain of $+1$, and has almost rail-to-rail performance at a single supply voltage of $3\,\text{V}$ (typical output range is $75\,\text{mV}$ to $2.93\,\text{V}$ driving $150\,\Omega$).

Figure placeholder

Figure 4.9: Effect of having the op-amp output idle at 0 V

Chapter 5

POWER SYSTEM DESIGN

This Chapter discusses how the delay line is powered, which is something which has required some consideration. As discussed in Specification point 2, the delay line must be powered from the input signal, and this method of providing power via this signal must require minimal modifications to the driving circuitry.

A naive solution would attempt to harvest the power of the input pulses themselves, and store power from each pulse to power the circuitry when the line is inactive. This is plausible because a 20 V peak-to-peak input signal, delivers 1.5 W of power into the delay line for the time it is active, as calculated by Equation 5.1.

$$P = \frac{V^2}{R} = \frac{10^2}{68} = 1.5\,\text{W} \tag{5.1}$$

The problem with this is approach is that one cannot guarantee that pulses will be delivered to the delay line with any frequency. When the store is storing all zeros, non pulses will be delivered to the delay line at all. This means that the delay line would have to power on with a predictable latency such that the first pulse could be delayed by the correct amount. Unfortunately the unpredictable delay of the the supply rails reaching the correct level, coupled with the hard to predict delay of the FPGA configuring from flash memory would break the timing budget.

Clearly a different solution to this problem is required, and this is what will be discussed in this Chapter. Section 5.1 will discuss the initially proposed solution which turned out to be flawed after further investigation, and Section 5.2 discussed a slightly more complex solution that overcomes the shortcomings of the former.

5.1 DC Offset Solution

Since the output of the valve circuitry is AC coupled **[Include schematic portion here?]** , a simple solution to interface with EDSAC is to add a DC offset to the output. This could be achieved as simply as adding a potential divider to the output of the store regeneration unit, stiff enough to provide a stable voltage for the delay line. This solution is illustrated in Figure 5.1.

This is an ideal solution in theory

## 5.2 Valve Power Supply Solution

5.2.1 SPICE Simulation

Figure 5.1: Proposed DC offset solution

Chapter 6

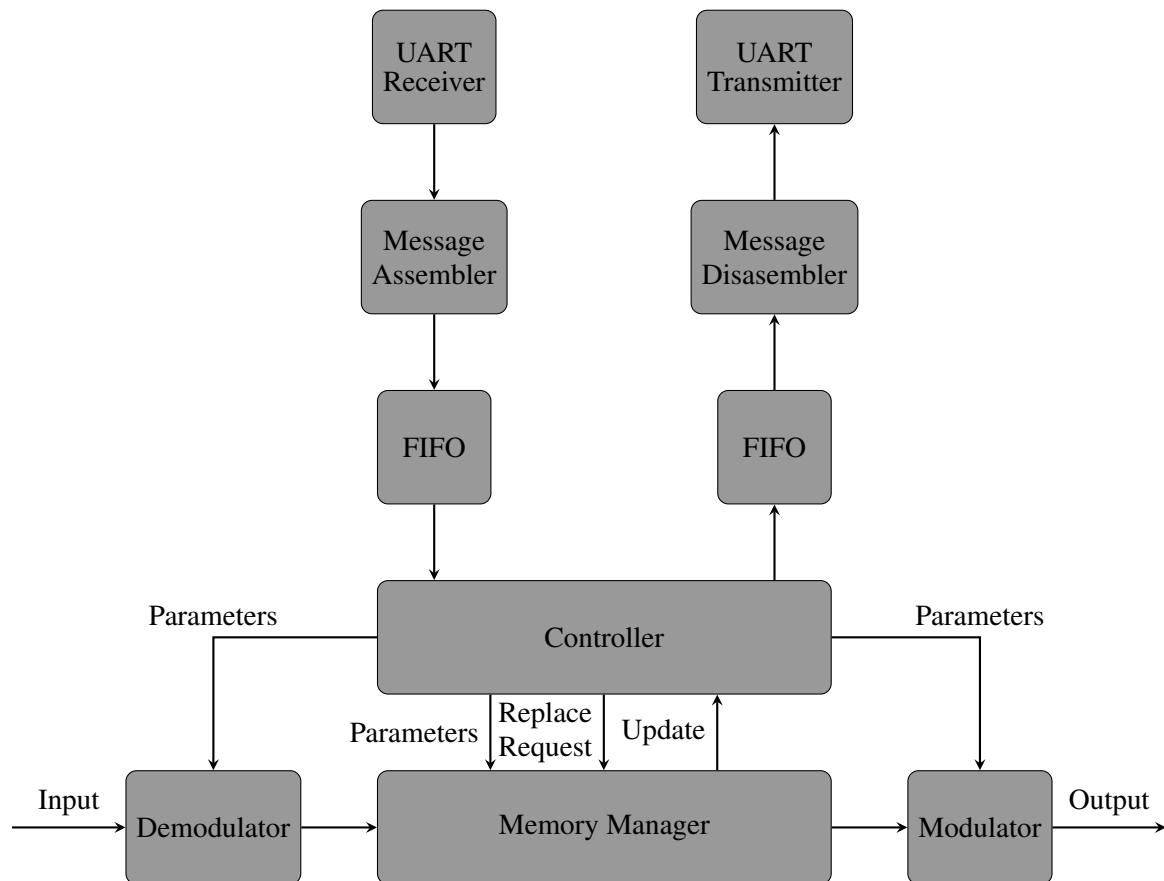# TEST HARNESS DESIGN

6.1 Digital Design

Figure 6.1: Proposed Architecture for Test Harness

6.1.1 HDL Design

6.1.2 HDL Verification

6.2 Analogue Design

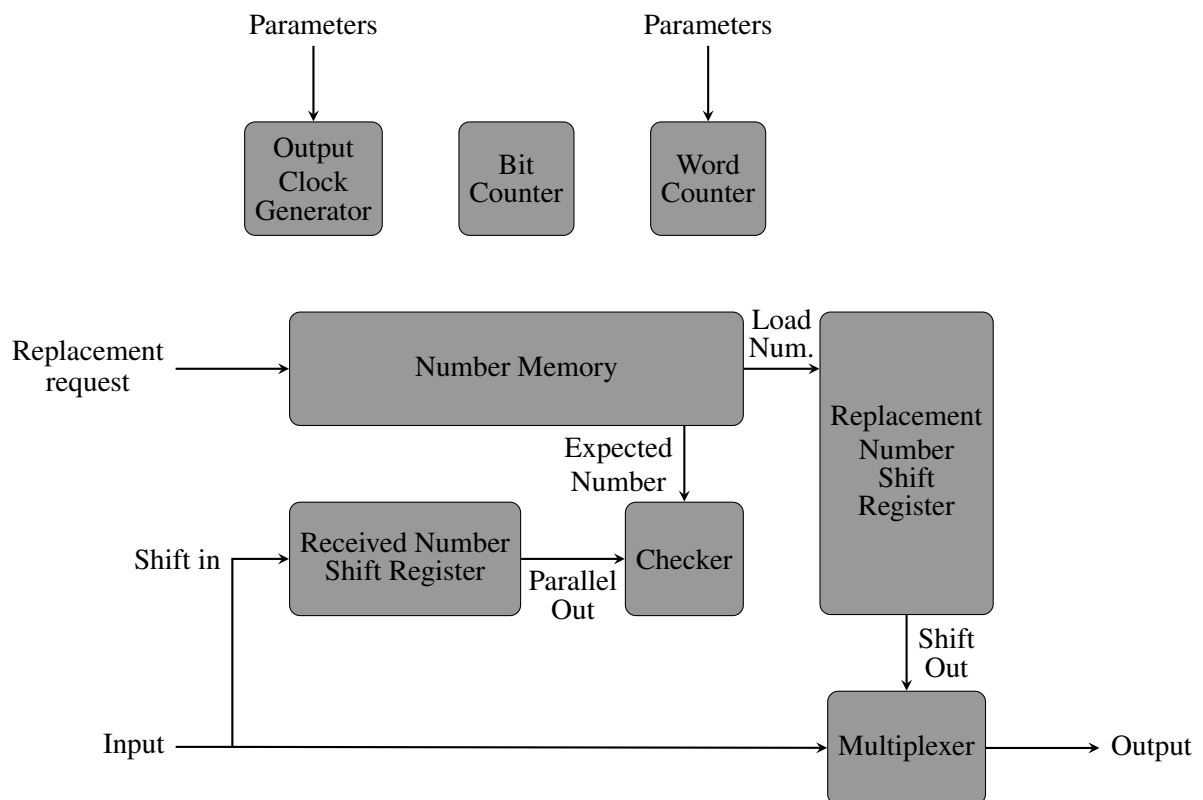### 6.2.1 Input

6.2.2 Output

6.3 Software Design

Figure 6.2: Proposed Architecture for Memory Manager

Chapter 7

SYSTEM INTEGRATION

7.1 Mechanical

7.2 Issues

7.2.1 FPGA Reset Issue

7.2.2 Signal Breakthrough

7.3 Measurements

Chapter 8

PROJECT PLANNING

Chapter 9

FINAL COMMENTS

Bibliography

[1] University of Cambridge. (2011, feb) EDSAC I, W.Renwick, M.Wilkes. Copyright Computer Laboratory, University of Cambridge. Reproduced by permission. Cropped from original photo. [Online]. Available: http://www.cl.cam.ac.uk/relics/jpegs/edsac_wilkes.jpg

[2] J. Tyler, "Mercury delay line emultion: Reconstructing memory for edsac," may 2017, this document is an unpublished research review written for the project.

[3] University of Cambridge. (2011, feb) M.V.Wilkes, mercury delay line (2). Copyright Computer Laboratory, University of Cambridge. Reproduced by permission. [Online]. Available: http://www.cl.cam.ac.uk/relics/jpegs/delay_lines.jpg

[4] M. Ward. (2011, jan) Pioneering edsac computer to be built at bletchley park. BBC News. [Online]. Available: http://www.bbc.co.uk/news/technology-12181153

[5] University of Cambridge, "Pioneer computer to be rebuilt," *CAM*, vol. 62, p. 5, mar 2011.

[6] F. da Cruz. (2013, nov) Programming the ENIAC. Columbia University. [Online]. Available: http://www.columbia.edu/cu/computinghistory/eniac.html

[7] K. S. Jones. (2001, dec) A brief informal history of the computer laboratory. University of Cambridge. [Online]. Available: http://www.cl.cam.ac.uk/events/EDSAC99/history.html

[8] J. Bonne. (2007, march) Transistor transition began in Allentown. The Morning Call. [Online]. Available: http://articles.mcall.com/2007-03-04/features/ 3712894_1_transistor-production-line-bell-labs-allentown-plant

[9] M. V. Wilkes and W. Renwick, "An ultrasonic memory unit for the EDSAC," *Electronic Engineering*, vol. 20, pp. 208–213, jul 1948.

[10] Farnell. (2017, aug) ICE40HX1K-TQ144 -FPGA, iCE40, PLL, 95 I/O's, 133 MHz, 1.14 V to 1.26 V, TQFP-144. [Online]. Available: http://uk.farnell.com/lattice-semiconductor/ ice40hx1k-tq144/fpga-1280-luts-1-2v-hx-144tqfp/dp/2362849

[11] *iCE40™LP/HX Family Data Sheet*, 3rd ed., Lattice, mar 2017.

[12] *iCEstick Evaluation Kit*, 1st ed., Lattice, aug 2013.

[13] C. Wolf and M. Lasser, "Project icestorm," http://www.clifford.at/icestorm/.

[14] *iCEcube2 User Guide*, Lattice, feb 2017.

[15] Vishay, *Small Signal Schottky Diodes, Single and Dual*, 1st ed., feb 2013.

[16] Diodes Incorporated, *Single Schmitt-Trigger Buffer*, 4th ed., mar 2014.

[17] Texas Instruments, *LMH6639 190MHz Rail-to-Rail Output Amplifier with Disable*, mar 2013.

Appendix A

Folder Structure

Foo

Appendix B
HDL Code Overview

[Is this section really necessary]

Appendix C

Schematics